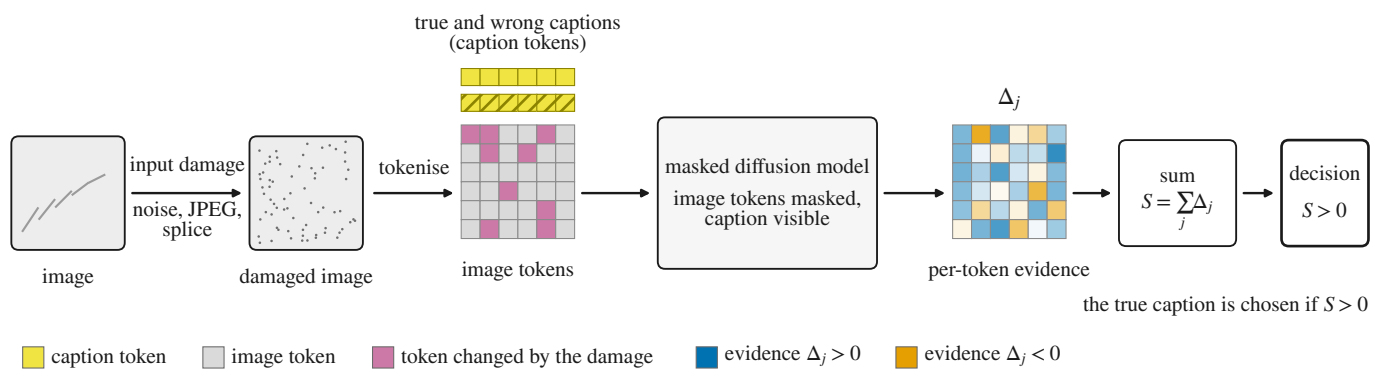
true and wrong captions
(caption tokens)
Δ_j
input damage
noise, JPEG,
splice
tokenise
masked diffusion model
image tokens masked,
caption visible
sum
$S = \sum_j \Delta_j$
decision
$S > 0$
image
damaged image
image tokens
per-token evidence
the true caption is chosen if $S > 0$
caption token
image token
token changed by the damage
evidence $\Delta_j > 0$
evidence $\Delta_j < 0$